

A GIS approach to modelling CO₂ emissions associated with the pupil-school commute

Alex Singleton*

Department of Geography and Planning, University of Liverpool, Liverpool, UK

(Received 12 November 2012; final version received 22 July 2013)

Concerns have been raised in numerous countries over declining rates of active transport to school. In a UK context, the pupil-school commute is estimated to contribute around 658 kilotonnes of CO₂ per year; however, tackling this issue effectively requires an improved understanding of how emissions can be modelled and mapped over a variety of scales. This paper implements a new estimation technique for the modelling of CO₂ emissions linked with the school commute that integrates both transport network-level routing and geographically disaggregate vehicle emissions data. The model is then applied to a national cohort of pupils in England. Areas demonstrating the highest emissions were typically more rural and/or comprising more affluent resident populations. Emissions were also shown to increase with school year, with larger step changes between educational stages reflecting the different geography of school locations. Furthermore, where secondary school entry policies were selective or based on a religious domination, average emissions were typically higher than in non-selective schools.

Keywords: GIS; CO₂ emissions; schools; vehicle routing

Introduction

Internationally, the rates of travelling to school by active transport (e.g. cycling or walking) are in decline (Tudor-Locke et al. 2001, Schlossberg et al. 2006, McMillan 2007, Trang et al. 2012), and the corollary switch to less sustainable modes of travel have been linked to negative effects on the environment in terms of increased emissions (Van Ristell et al. 2012), increasing traffic congestion around schools (Collins and Kearns 2001) and health impacts related to lower physical activity levels (Faulkner et al. 2009, Meron et al. 2011) or pollutant exposure (McConnell et al. 2010).

In a UK context, schools account for 15% of total public sector emissions (DCFS 2010), which in England are estimated to be the equivalent of around 9.4 million tonnes of CO₂ per year (SDC 2006). Seven per cent (658 kilotonnes) of this total is associated with the pupil-school commute, and as such, there are significant environmental benefits for pupils to adopt more sustainable travel behaviour.

International research on commuting to school reveals that mode choice is impacted by multiple interacting factors including actual and perceived distance to the school (McDonald 2007, Müller et al. 2008, Lang et al. 2011), road infrastructure (Ewing et al. 2004, Bejleri et al. 2010), urban form (McMillan 2007, Mitra et al. 2010, Panter et al.

*Email: alex.singleton@liverpool.ac.uk

2010, Cui *et al.* 2011, Broberg *et al.* 2013), ethnicity (McDonald 2007), socio-economic status (Wilson *et al.* 2010, Roth *et al.* 2012, Su *et al.* 2013) and lifestyle factors (Babey *et al.* 2009). It has also been argued that policies enabling school choice can lead to longer commuting distances, which has a corresponding impact on mode choice and emission levels (Marshall *et al.* 2010, Wilson *et al.* 2010, Van Ristell *et al.* 2012). The use of geographic information systems (GIS) as a framework for analysis is prevalent in many studies in this area. Examples include calibrating environmental measures related to active transport (Panter *et al.* 2010, Dalton *et al.* 2011, Wong *et al.* 2011) including accessibility and walkability (Bejleri *et al.* 2010, Giles-Corti *et al.* 2011); estimating probable routes taken (Stigell and Schantz 2011, Larsen *et al.* 2012) or energy expended (Cooper *et al.* 2010); and finally, assessing pollutant exposure along routes (Ashmore *et al.* 2000).

The mode of transport adopted by pupils for their commute to school has obvious implications for CO₂ emissions. With a few recent exceptions (Wilson *et al.* 2007, Marshall *et al.* 2010, Van Ristell *et al.* 2012), there have been limited studies to date that estimate the emissions impact of these journeys, with only Van Ristell *et al.* (2012) examining this at a national scale for England (albeit using a sample). In two of these studies (Marshall *et al.* 2010, Van Ristell *et al.* 2012), different types of regression analysis were used to predict mode choice given a range of influencing factors (Wilson *et al.* 2010). From these models, CO₂ emissions are estimated by multiplying the distance and frequency of trips between locations by the average CO₂ attributed to specific modes of transport. Commuting distance is treated differently by these studies, with Van Ristell *et al.* (2012) adopting straight-line distance, whereas Marshall *et al.* (2010) and Wilson *et al.* (2007) implemented network-based distances derived using shortest or actual commuting paths. The advantage of straight-line distance is that it can be computed quickly; however, this will typically underestimate the true distance travelled, given that actual routes follow street/rail/footpath topology. However, the estimation of shortest or quickest path along a network is computationally intensive, thus creating an additional challenge for studies concerning either a very large sample or total population survey. A second consideration is what CO₂ (g/km) value should be attributed to different mode choices? The use of national averages is prevalent in most existing models, however, may overestimate or underestimate geographical differences in actual emission characteristics; for example, relating to issues such as the mix of cars typically owned within different types of area.

As such, there is a challenge to create transport-linked CO₂ emission estimates that better account for geography; both in the estimation of the distances travelled over multiple transport networks, and where possible, through the integration of mode-choice emission input parameters that are sensitive to geographical context. This complex task is computationally intensive, however, as illustrated in this paper, can be enabled through the application of geocomputation (Openshaw and Abrahart 2000). A geographically sensitive model is specified that integrates network-level (road/rail) routing with geographically disaggregate vehicle emissions data to provide estimates of pupil school commute linked CO₂ emissions at an individual scale, and for a national cohort of pupils. These results are compared to those achieved using a simple model that is calibrated with straight-line distance and national average CO₂ estimates. The paper concludes with an exploration of the socio-spatial structure of the geographically sensitive estimates.

A geographically sensitive model of CO₂ emission estimates

An estimate of CO₂ emissions attributed to a pupil (p) school commute (k_p) is defined in Equation (1) where d is the distance travelled between origin i (pupil domicile unit

postcode) and destination j (school unit postcode) utilising transport mode t with associated CO_2 (g/km) value e . Furthermore, in the case of transport by car or light rail, these average vehicle emissions can also vary by location g . This parameter is used flexibly within the model, with location being specified as either a lower super output area (LSOA: 400–1200 households) in the case of car travel, or for light rail, the nearest network to the pupil domicile. The CO_2 (g/km) estimates that are used to calibrate the model are discussed later and detailed in [Figure 1](#) and [Table 1](#). A final weighting parameter w is required to adjust an estimate where average occupancy is not already accounted for in the emission value associated with the transport mode choice. In this model, it only includes those pupils whose travel mode is by car sharing, where a weighting of 0.5 is presumed (i.e. an occupancy of two).

$$k_p = 2 \left(d(i_p, j_p, t_p) e(t_p, g_p) w(t_p) \right) \quad (1)$$

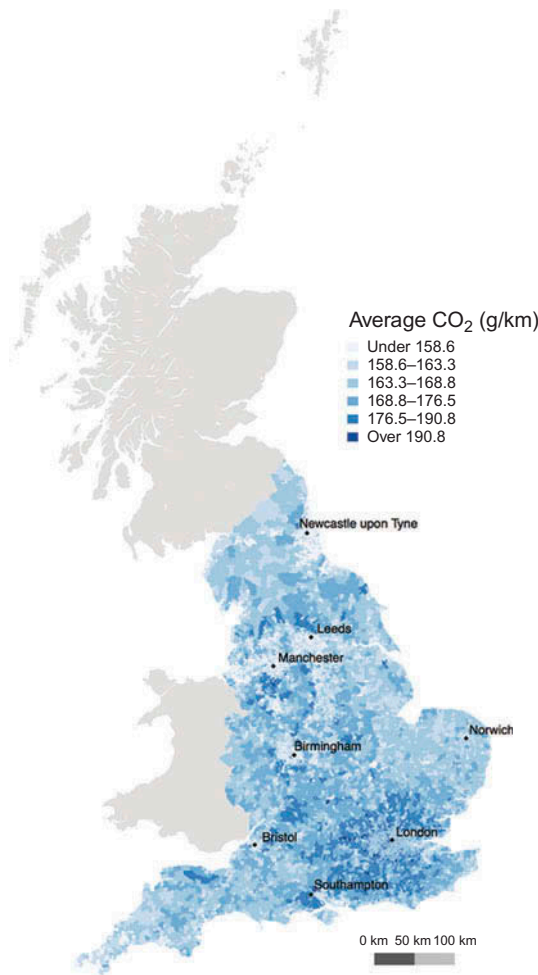


Figure 1. Average CO_2 values, given in g/km, for those cars registered within lower-layer super output areas.

Table 1. CO₂ emissions for different transport modes (Coley 2002, DEFRA 2011, Tranter 2012).

Transport mode	Average CO ₂ (g/km/passenger)	Average occupancy
Taxi	150.3	1.4
Bus (London)	85.7	16.7
Bus (non-London)	184.3	6.3
Coach	30.0	16.2
Light rail – average	71	–
London (DLR)	68.3	93
Birmingham/Midlands	70.5	36
Newcastle	103.0	53
Croydon	44.3	55
Manchester	39.5	59
Nottingham	#	36
Sheffield	96.8	40
National rail	53.4	–
London underground	73.1	–
Cycling	8.3	NA
Walking	11.4	NA

Notes: # = No network-specific estimate was available; the model was implemented using the light rail average as surrogate.

NA, not applicable.

The model is empirically driven, requiring a number input of data sets. The Department for Education (DfE) supplied a national cohort data set containing individual pupil-level attributes such as age, gender, ethnicity and home postcode, with a coverage for all pupils in state-maintained schools in England, from nursery through to post-compulsory education.¹ These data are collected by the DfE for policy and administrative purposes, and it is mandatory for schools to complete a return, with three separate surveys completed throughout the year, and spaced between the spring, summer and autumn terms. Data are available going back to 2002; however, between 2007 and 2012, there was also an additional requirement for schools to submit details of ‘usual’ transport mode choices adopted by pupils for their journeys to and from school. This represents the best estimate of actual mode choice, however, it should be noted that students may adopt different modes of travel at the start and end of the day (Schlossberg et al. 2006, Southward et al. 2012); however, in a UK context, travel mode to and from school has been shown to correlate strongly (van Sluijs et al. 2009). The DfE supplied access to the Spring 2011 data, which at the time of request was the latest available release. The total number of records in this pupil database was 7,532,205; however, from these, 158,700 were excluded as they either did not have a domicile postcode, or the postcode provided was invalid; mode choice was null, specified as a ‘border’ (attending a boarding school) or ‘other’ (an uncategorised mode choice), thus leaving a final data set of 7,373,505 observations.

The second key data set integrated into the model contains the average of CO₂ (g/km) emitted by those cars that were registered to addresses within each LSOA. Averages were created for this project by the Department for Transport (DfT) as an aggregation of Driver Vehicle Licensing Agency (DVLA) data. This records an associated CO₂ (g/km) emission value for those vehicles that are registered to specific addresses. These emission values are derived using a rolling road test as set out in the European legislation directive 91/441/EC and are a component of emissions approval. However, when these values are aggregated at an area level, they may underestimate true emissions averages, given that

the cars registered after March 2001 are only included in these tests. For example, the average vehicle CO₂ (g/km) from the DVLA data is 162.9, which differs from the Department for Environment, Food and Rural Affairs (DEFRA) national estimate of 183.2 (DEFRA 2011). However, without actual data on those cars driven to school, these estimates represent a plausible alternative. If the LSOA-level average car emission rates are mapped (see Figure 1), there are clear geographic variations; for example, highlighting how cars registered in the areas surrounding London and the south of Manchester have on average higher emission rates. Exploring the underlying causes of these variations lies outside the scope of this paper; however, this map does illustrate the importance of including geographically contextualised inputs into models when estimating emission levels at a disaggregate spatial scale.

The model also requires average vehicle CO₂ per kilometre emission estimates for non-car transport. These are given in Table 1 and are sourced from a number of publications (Coley 2002, DEFRA 2011, Tranter 2012). In general, these estimates are calculated as an interaction between the average emission characteristics of a type of vehicle and the typical occupancy, which is why for a single mode (e.g. light rail) there are multiple values for those schemes operating in different parts of England. The rates for vehicles were all derived from DEFRA (2011) and are the official (and annually updated) figures used in government statistics concerning emissions.

An innovative feature of the model is to estimate the distances pupils travel to school along both the street and rail networks. There are numerous ways in which these network distances could be estimated, ranging from online services to various open-source or commercial software. All implementations involved the automated supply of origin destination locations and the calculation of distances over either the quickest or shortest path along the transport network. For automotive journeys, a wide variety of solutions are available to calculate routes along the road network, and those tested included the application programming interfaces (API) for Google Maps² and Bing Maps,³ alongside the open-source routing library Routino⁴ that utilises OpenStreetMap⁵ data. The API for both Google Maps and Bing Maps enabled a query to be sent to their respective online servers, a route calculated and returned in a standardised format. Routino is an open-source software that can be compiled for both Linux (<http://www.linux.org>) or OS X (<http://www.apple.com/uk/osx/>) and installed locally. A query can be sent to the library using the terminal, and will return a result that calculates a route over OpenStreetMap street network data. Other options considered, but not tested, included various online routing services that also integrated with OpenStreetMap data such as OpenRouteService⁶ or YOURS.⁷ Given the non-commercial nature of these services, the impact on their servers and bandwidth of calculating such large volumes of routes would have been significant; and as such, these services were not tested. Both the Google and Bing APIs were subject to search limits. For Bing, this was 125,000 sessions or 500,000 transactions within a 12-month period (developers), and for Google, 2500 per 24-hour period. However, in addition to these limits, both services appeared subject to adaptive firewall technology, thus preventing multiple queries in succession. The calls made were within the bounds of the daily usage limits, however, at various random intervals, would fail to return a result. The reasons for these failures were not clear, and indeed results could be returned if the failed origin destination queries were re-run at a later stage. A further limitation to the Google service was a requirement in their terms of use that any results derived from routing queries have to be displayed on a Google map after the trip calculation. Thus, although the distance could be extracted, it was also a necessity to display the returned route on a map, which has large performance implications, considerably slowing down batch calculations given the rendering time of

maps would comprise an unnecessary processing overhead. Overall, both Google and Bing services provided a reasonable solution for calculating routes, however, daily query limits, in addition to adaptive firewall issues limited their utility when calculating routes in large volumes. Furthermore, from a methodological perspective, these services are closed source, and there is little information on those algorithms utilised or metadata about the network over which the routes are calculated. An alternate solution evaluated was the open-source Routino libraries that can run under Linux or OS X. The Routino libraries (version 2.2) were compiled from source under OS X Lion (<http://www.apple.com/uk/osx/>), and an extract of the Great Britain OpenStreetMap data was imported. The default configuration of the Routino library was implemented with the exception of enabling pedestrian access to 'trunk roads' (larger, often dual carriageways). This corrected problems where pedestrian access had not always been explicitly labelled in the underlying OpenStreetMap data, and because of this, the routing algorithm would occasionally calculate implausible and very long routes to avoid the 'trunk roads'. Solutions were compared by examining a set of 500 randomly selected origin–destination pairs from the pupil database. There was a strong positive linear relationship between the distance estimates with Pearson correlation coefficients ranging between 0.98 and 0.99 ($P < 0.001$). As such, the routing solutions are reasonably interchangeable; however, Routino was implemented in the model as this offered the lowest computational processing time given that the calculations were completed locally, and furthermore, unlike the two online systems is open source, thus enabling greater transparency of the results.

For non-road journeys (rail, underground, light rail), appropriate online routing applications were not available, as although it is currently possible for services to supply quickest travel times, actual track distances are not provided. In the case of the Google Maps API, some railway routes have started to be added (at least visually); however, these are not uniformly available. As such, a custom solution was developed by importing rail network data into a PostgreSQL⁸ database with the PostGIS⁹ (spatial enabling) and pgRouting¹⁰ (routing enabling) extensions. Railway network and station location data were extracted from national mapping agency open data (Ordnance Survey: Meridian 2¹¹); and for light rail and the London Tube, these were exported from OpenStreetMap. The topologies of these data were extensively checked to ensure that routes could be correctly calculated after they were imported into the database. For pupils travelling on non-road transport, a nearest neighbour calculation was used to return the closest station to both the pupil domicile, and additionally the school attended, with routes being calculated between these two station locations. Rail has reasonably universal coverage; however, for both tube and light rail, these modes of transport are only found within certain areas. As such, an error check was required to examine the plausibility of the assignment of both pupils and schools to a nearest station (see Equation (2)). For each distance (D) between a pupil (p) domicile and their nearest (x) station, a Tukey outlier filter (Tukey 1977) was calculated using the quartile (Q) and inter-quartile range (IQR) specific to each network (n).

$$[(Q1_n - 1.5IQR_n) < D_{xp} < (Q3_n + 1.5IQR_n)] \quad (2)$$

An outlier check is necessary given that a light rail, tube or rail mode choice could have been specified in error, which would result in pupils being assigned to a station at an unrealistic distance from their domicile. The same calculation was implemented for schools, thus supplementing p in the above equation with s for school. Those distance scores returned as outliers were removed from the analysis, as were those pupil journeys that returned a valid result but were equidistant between two networks. For these cases, there was no way of

determining which network was used, and was an issue specific to London, where two light rail networks are present within reasonably close proximity, however, are not physically connected. These procedures were not implemented for the railway network given that station coverage is national, and as such, based on a nearest neighbour assignment, it would be difficult to pick out those matches that were in error. However, the distances calculated from erroneous origin destination pairs would be detectable as outliers (i.e. travelling implausible distances), which is discussed later in the context of all transport mode choices. Out of 16,048 light rail and tube journeys, 1621 were identified as outliers ($\sim 10\%$). Finally, a further error check compared the distance to the nearest station with the distance to the school attended. If the straight-line distance to the school was shorter than the distance to the station, then the mode was presumed to be in error, and as such, these pupils were also excluded from the analysis. For all non-road journeys, this removed another 1914 records.

When implementing distance calculations, either the shortest or quickest journeys were estimated. For road network calculations, these were implemented as the quickest routes, with access and speed restrictions appropriate to the mode of transport. For example, a car would not be able to access pedestrian footpaths running through a park. Although bus routes are beginning to become available for some operators in England, coverage is currently limited. Furthermore, dedicated school bus routes are not known. As such, these distances were also calculated as road optimised quickest paths. For rail-based calculations, the shortest path was calculated, optimising routes based on the track length.

After the model input data were assembled into a PostgreSQL database and routing solutions were created for both rail and road journeys, these were coupled using the statistical programming language R (<http://www.r-project.org>). A pseudo-code implementation of the model is detailed below:

- (1) For each pupil, identify the mode choice
 - (a) Implement a rail- or road-based distance calculation appropriate to the mode
 - (b) Store the distance travelled
- (2) Check that the pupil distances returned are plausible, and extract outliers or detected errors
- (3) For each valid pupil commute
 - (a) Retrieve an appropriate geographically sensitive CO₂ (g/km) emissions estimate depending on location and mode choice
 - (b) Calculate CO₂ (g/km) for each pupil trip
 - (c) If necessary, weight CO₂ (g/km) to account for a car share arrangement
 - (d) Store emissions estimate

At stage 1(a) in the above pseudo-code, distances are returned using either road or rail network distances. All rail routes were returned, however, the Routino-based road network calculations could not return a distance for 24,319 routes despite having valid co-ordinates specified for both origins and destinations. These errors occur because the algorithm could not find a proximal node within the OpenStreetMap network data and can occur where the network coverage is lacking.

After calculating appropriate network distances, an error test was implemented to check the plausibility of the results. There are a number of reasons why errors may occur, including

- that the domicile postcode is valid but specified in error
- that the school postcode may be an error (although unlikely given the administrative importance of these data being correct)

- that there has been a geocoding error
- that the routing calculation had returned an implausible result (possible for road calculations that are dependent on underlying OpenStreetMap (OSM) data quality)

A modified Tukey outlier filter (Tukey 1977) was again implemented to examine each pupil commute (K_p) against the quartile (Q) and inter-quartile range (IQR) for all distances disaggregated by the combination of year group (a), local authority (g) and mode type (m). This was necessary to account for geographic differences in transport infrastructure, differences in those distances typically travelled using different types of transport, and finally, to control for how travel behaviour changes between age groups. This was implemented using Equation (3):

$$[(Q1_{mgq} - 1.5IQR_{mga}) < K_p < (Q3_{mga} + 1.5IQR_{mga})] \quad (3)$$

This analysis resulted in a final database of 6,912,226 records being returned for input into the CO₂ estimation model. The distances were then multiplied by an appropriate CO₂, given in g/km, for the allocated mode choice, and if necessary, weighted, as described earlier in this section, and then stored back in the database.

The validity of the model can be explored through comparison with official national estimates, which as described at the start of this paper are 658,000 tonnes of CO₂ per year for the UK. If the emission value for each pupil is multiplied by 5 (days of the week) and then by 38 (weeks of schooling), these scores give a summed value for England of 618,602 tonnes CO₂ per year. These are reasonably in line with UK statistics, given that they exclude Northern Ireland, Wales and Scotland; and although there are no methodological details of how the official UK estimates were calculated, these would most probably utilise Euclidian distances, which as discussed earlier, would systematically underestimate the real distances travelled. It should also be noted that the estimates calculated in this paper exclude those pupils in non-state schools and those records identified as outliers. Further validation of the model was implemented by comparing the results of the geographically sensitive model with an alternate simple model that supplemented transport network distance with straight-line distance, and swapped the location specific vehicle CO₂ emission values (previously given in Figure 1 and Table 1) for national averages (DEFRA 2011). A map illustrating the differences between the CO₂ emissions estimated by the geographically sensitive and the simple models are shown in Figure 2. The choropleth units relate to local authority boundaries, which are an administrative geography for England. Positive scores within these zones indicate that the simple model is overestimating emissions relative to the geographically sensitive model, with negative results illustrating the inverse pattern. The top ten local authorities with the highest overestimated and underestimated values shown in Table 2. These patterns demonstrate a degree of systematic error, with the simple model more prevalently exhibiting underestimated values for urban areas and the opposite pattern observed for rural local authorities. The implication of these differences reinforces a recommendation for models to better integrate local geography, and is arguably acutely important if the results are to be presented at sub-national scales.

The geography of CO₂ emissions linked to the school commute

The model presented in the previous section created an estimated CO₂ emission value for each valid pupil-school commute. In this following section, these estimates are explored

by various factors including school characteristics and year groupings, pupil locations, and a geodemographic classification that provides a summary of multiple characteristics about the areas in which the pupils are domiciled. These descriptive analyses illustrate key characteristics of the English national geography of CO₂ emissions linked to the school commute.

The national geography of CO₂ emissions was explored by calculating the pupil average daily CO₂ emission values within local authority administrative areas (see Figure 3).

The local authorities with the highest values attributed to the school commute were Sevenoaks with estimated average daily emissions of 1184 g of CO₂; Canterbury 1115 g; Tonbridge and Malling 1110 g, East Lindsey 1086 g and South Holland 969 g. Unpicking this geography is complex given the national coverage, however, in general, those areas with higher average emissions appeared to be in lower density and more rural areas. This geography can be explored further by profiling the model results by a geodemographic classification linked to each pupil postcode. The importance of contextual effects




Figure 3. Aggregated average daily emissions associated with the school commute.

on children's school travel mode choice has been explored elsewhere (Sidharthan et al. 2011). Although some have criticised their use (Goss 2003), there is a diverse and international legacy of successful application of geodemographics in the exploration of urban socio-spatial structure; with utility demonstrated across a wide variety of public sector domains (Singleton and Spielman 2013). The ACORN geodemographic classification from CACI Ltd (<http://www.caci.co.uk/>) was made available for this analysis and comprised three hierarchically nested levels of 5 categories, 17 groups and 56 types. Each of the clusters within the typology is named and described to give end-users an impression of those salient characterisers of the classified areas. The average daily CO₂ (g) associated with the pupil-school commutes are presented by the ACORN typology in Table 3.

Table 3. Average CO₂ (g) emitted during the pupil-school commute disaggregated by the ACORN geodemographic typology.

Categories	Groups	Types	Average CO ₂ (g)
Wealthy achievers	(A) Wealthy executives	(1) Wealthy mature professionals, large houses	734.7
		(2) Wealthy working families with mortgages	608.4
		(3) Villages with wealthy commuters	844.2
		(4) Well-off managers, larger houses	743.5
	(B) Affluent greys	(5) Older affluent professionals	732.5
		(6) Farming communities	1231.4
		(7) Old people, detached homes	672.2
		(8) Mature couples, smaller detached homes	668.8
	(C) Flourishing families	(9) Older families, prosperous suburbs	497
		(10) Well-off working families with mortgages	579.2
		(11) Well-off managers, detached houses	655
		(12) Large families and houses in rural areas	791.3
Urban prosperity	(D) Prosperous professionals	(13) Well-off professionals, larger houses and converted flats	315.8
		(14) Older professionals in suburban houses and apartments	454.4
		(15) Affluent urban professionals, flats	302.1
	(E) Educated urbanites	(16) Prosperous young professionals, flats	303.4
		(17) Young educated workers, flats	422.1
		(18) Multi-ethnic young, converted flats	315.5
		(19) Suburban privately renting professionals	444.1
	(F) Aspiring singles	(20) Student flats and cosmopolitan sharers	381.1
		(21) Singles and sharers, multi-ethnic areas	364.2
		(22) Low income singles, small rented flats	420.6
		(23) Student terraces	481.2
	Comfortably off	(G) Starting out	(24) Young couples, flats and terraces
(25) White-collar singles/sharers, terraces			450.7
(H) Secure families		(26) Younger white-collar couples with mortgages	551
		(27) Middle income, home owning areas	610.5
(I) Settled suburbia		(28) Working families with mortgages	457.5
		(29) Mature families in suburban semis	444.6
		(30) Established home owning workers	423.5
		(31) Home owning Asian family areas	454.2
		(32) Retired home owners	573.3
		(33) Middle income, older couples	467.8
		(34) Lower incomes, older people, semis	492.2

(Continued)

Table 3. (Continued).

Categories	Groups	Types	Average CO ₂ (g)
Moderate means	(J) Prudent pensioners	(35) Elderly singles, purpose built flats	496.7
		(36) Older people, flats	472.3
	(K) Asian communities	(37) Crowded Asian terraces	310.9
		(38) Low income Asian families	322.2
	(L) Post-industrial families	(39) Skilled older families, terraces	402.1
		(40) Young working families	386.6
Hard-pressed	(M) Blue-collar roots	(41) Skilled workers, semis and terraces	419.7
		(42) Home-owning families, terraces	339.3
		(43) Older people, rented terraces	364.1
		(44) Low-income larger families, semis	396.6
	(N) Struggling families	(45) Low income, older people, smaller semis	382.2
		(46) Low income, routine jobs, terraces and flats	382.3
		(47) Low-income families, terraced estates	353.4
		(48) Families and single parents, semis and terraces	329.5
		(49) Large families and single parents, many children	353.2
		(50) Single elderly people, council flats	372.5
	(O) Burdened singles	(51) Single parents and pensioners, council terraces	358.5
		(52) Families and single parents, council flats	345
	(P) High-rise hardship	(53) Old people, many high-rise flats	380.4
		(54) Singles and single parents, high-rise estates	349.9
	(Q) Inner city adversity	(55) Multi-ethnic purpose built estates	283.5
		(56) Multi-ethnic crowded flats	279.1

A number of interesting patterns are highlighted. Firstly, reinforcing those previous observations from Figure 4, some of those geodemographic types with the highest emission values are typically found within more rural areas, for example, 'Villages with wealthy commuters', 'Farming communities' and 'Large families and houses in rural areas'. Secondly, the most affluent types typically have higher emission values (e.g. those within the category 'Wealthy achievers'). The inverse of these patterns emerges in those types that pertain to areas where residential populations have more constrained finances and are typically found within urban areas (e.g. within the categories 'Moderate means' and 'Hard-pressed'). Drivers of these patterns will likely relate to both the characteristics of the built environment in the places where the geo-demographic clusters are typically found, and also the behavioural choices of their typical resident populations. For example, affluent populations may more prevalently choose to drive over other forms of transport; or, those cars purchased by people living within more rural areas may bias vehicles with larger engines that are more suitable to variable terrain. Exploring the underlying motivations for such choices lies outside the scope of this study; however, to illustrate the influence that location and transport mode choice has on CO₂ emissions between geodemographic types, a scatter plot of commuting distance and the percentage of pupils adopting active transport is shown in Figure 4.




Figure 4. The geodemographics of CO₂ emissions, commuting distance and active transport.

The points plotted in Figure 4 are ACORN types, coloured by their nested category and are scaled by the average CO₂ (g) emitted, with smaller points equating to lower emissions. As might be expected, there is a strong negative linear relationship between the distance travelled to school and the adoption of active transport, with pupils resident in areas classified by those types travelling further, typically contributing more CO₂ on average per commute. The disaggregation of this distribution by geo-demographic categories is quite striking, illustrating again that those types containing more affluent populations are on average contributing higher levels of CO₂ emissions per daily commute, travelling further, and with lower rates of active transport.

Preferences for active transport will not be homogenous across school years given the variable geographic distribution of schools at various educational stages, and additionally, growing independence of pupils over time. Figure 5a presents a disaggregation by school year for the average CO₂ (g) emitted during the school commute. The coloured bars represent different educational stages from nursery (years N1, N2; aged 4–5), through reception (year R; aged 5), and then into primary school (years 1–6/7; aged 6–11/12), secondary school (years 6/7–11; aged 11/12–16) and finally, post-compulsory education (years 12–13; aged 16–18). Average impact grows as school year progresses with step changes occurring as pupils transfer from nursery school to primary school, primary school




Figure 5. (a) Average CO₂ emissions by school year (b) percentage of pupils travelling using active transport by school year.

to secondary school, and finally, into post-compulsory education. The step changes relate to the different distribution of schools and/or travel options; for example, primary schools are typically smaller than secondary schools and have a more defuse geographic coverage. The gradual increase in emissions between years is also interesting; more so given that this pattern follows for primary, secondary, and post-compulsory education. The distances travelled to school demonstrated the same gradual increase, which is perhaps not that unexpected given their use in the estimation of the CO₂ emitted; however, these are also matched by an overall gradual decline in active transport (see Figure 5b). One issue with these observations is that the patterns observed are not those of a single cohort of students over the duration of their education, rather those pupils within schools during a single academic year; and as such, it is problematic inferring too strongly about the changing behaviour of individuals over time.

The availability and preferences for different types of school will also have an interacting effect on overall emission rates; and as such, the final attribute explored in this paper considers secondary school admissions criteria. Only secondary schools are considered in this section given that there are a greater diversity of admissions criteria, and that these will typically have greater impact on distances pupils travel to school given their more defuse geographic coverage. Within England, there are a number of different systems that govern admissions to secondary schools. These differ regionally, and broadly can be divided into selective (admissions based on a test) and comprehensive (admissions without a test). In those areas with only selective admissions, a further category exists ('Modern') which is applied to those schools where pupils would attend if they failed to attain a place at a selective school. Furthermore, some schools also have a requirement (variably enforced) that pupils should be from a specific religious denomination; providing a secondary and sometimes overlapping constraint on choice.

Parents choose which secondary school to send their children; however, the allocation of pupils into schools is bound by the availability of places. Higher performing schools will typically have increased demand, often leading to over subscription, where more applications are received than available places. Although various criteria are used to allocate pupils to available places in non-selective schools, a commonly used differentiating factor is straight-line distance between a pupil domicile and the school (Singleton et al. 2011).

Table 4. Emission characteristics disaggregated by secondary school selection criteria and presence or absence of religious denomination.

Type of school	Average CO ₂ (g/km)	Average distance (km)	Active transport (%)
No denomination	642.3	3.8	49.3
Religious denomination	871.9	4.8	30.5
Comprehensive	629.0	3.8	47.6
Modern	826.5	4.5	45.0
Selective	1529.6	7.7	19.5

As such, those pupils who are more proximal to the school, on average, will share a greater chance of admission, all other criteria being equal. For selective schools, these distance criteria will not typically be applicable, and where the performance of the school is high, this may attract pupils from more disparate locations. For non-selective schools, religious denomination can often usurp distance-based entrance criteria, and again, has the effect of increasing the average distances that pupils will travel to these schools. The impact of school policy on commuting-linked emissions and the distances pupils travel to school are summarised in Table 4. This shows that pupils attending religious schools are on average travelling further to school, with lower rates of active transport and higher average CO₂ in g/km travelled. A similar, but more extreme pattern is exhibited for those attending selective schools, and additionally, again, with a very low rate of active transport. These data illustrate that school admission criteria impact strongly on emissions and support previous work in this area (Van Ristell et al. 2012) that argue for more sustainable ‘neighbourhood schools’, where pupils are assigned to a nearest (or zoned) school, thus minimising travel and increasing the prevalence for active transport.

Conclusions

This paper has explored the geography of CO₂ emissions linked with the pupil-school commute in England. A method of estimation is presented that accounts for local geography through the inclusion of transport network distances and vehicle emission characteristics that are geographically disaggregate. This inclusion represents an incremental improvement on established methods that typically utilise straight-line distances and national average vehicle emission values. A second innovative feature of the model is that the calculations have been scaled for a near-total population of school pupils, thus creating CO₂ emission estimates for around 7 million individual commutes over the transport network. Previous work in this area has relied on samples (Van Ristell et al. 2012) or small focused case studies (Wilson et al. 2007, Marshall et al. 2010), and as such, this work represents the first time a model of such disaggregate scale and extensive coverage has been implemented in the context of pupil-school commuting. When comparing results to alternate estimates generated using straight-line distance and national average vehicle emission values, it was found that these simplified models had a tendency to underestimate emissions relative to a geographically sensitive model for a number of more urban local authorities. This likely relates to a composite of factors, however, an obvious impact would be the added complexity of the urban street network (density and directionality of traffic), which would impact road travel distances.

The results of the geographically sensitive model were mapped and then profiled by a series of categorical descriptors to elicit the types of areas in which the pupils were domiciled, alongside the characteristics of those schools attended. This analysis supported findings from the wider literature that has illustrated links between socio-economic status,

the built environment and travel demand/mode choice (Ewing and Cervero 2001, Handy et al. 2005, Brownstone and Golob 2009). Higher average emissions per pupil were more prevalently found in lower density areas and/or with more affluent populations. The importance of geographic context has been shown elsewhere to influence pupil mode choices (Sidharthan et al. 2011), and as such, there may also be a degree of self-reinforcement within or between proximal zones that have populations predisposed to adopting less sustainable methods of commuting to school. From a policy perspective, this may make 'neighbourhood' an appropriate scale at which to encourage more sustainable behaviour.

Average emissions increased with each year of academic progression, with larger step changes occurring between the different educational stages (e.g. primary to secondary). The type of school that parents choose for their children was also shown to have an impact on emissions. Pupils attending schools where entrance was governed by selective admission policies (based on a test) tended to exhibit a higher propensity for travelling longer distances, with a correspondingly lower rate of active transport and increased average emissions. Similar patterns were also observed for those pupils attending schools where religious denomination formed a component of the entrance criteria. It is argued that this relates to an absence of entrance criteria linked to proximity measures, which is a common metric used in non-selective schools where demand outstrips supply. From a policy perspective, this supports those arguments made elsewhere in the literature (Van Ristell et al. 2012) that significant potential benefits for reducing CO₂ emissions could be attained through policy that encourages schools with localised catchments.

The model presented in this paper has illustrated a method of improving geographical sensitivity in transport-linked emissions models; however, it has also highlighted some areas for extension. For example, actual bus routes could be integrated when national coverage of these data becomes available. This would also enable multi-mode routing to be considered, which may be especially important in areas such as London where a more integrated public transport infrastructure exists. Further model refinements could also be envisaged to better account for the physical characteristics of road routes, for example, by measuring terrain, assessing the probable number of stops at lights and accounting for other network characteristics such as speed limits. These data could be used to weight CO₂ estimates, as they would have an impact on car emissions. In addition to these methodological enhancements, a number of substantive directions for future research emerged from the empirical analysis. Firstly, to explore how emission patterns changed between 2007 and 2012, corresponding to when the national mode choice data were collected; and secondly, to explore how emissions and mode choices change between educational stages through the linking of individual pupil records over time.

This paper has described and then implemented an alternate way in which CO₂ emissions linked with the pupil-school commute can be estimated for a national cohort of pupils; building on previous work by encapsulating the topological structure of the transport network and geographic differences between average vehicle emission characteristics. The substantive findings are supportive of previous studies in this area, indicating that built structure, affluence and school admission policies all have impact on overall CO₂ emissions.

Acknowledgements

The author thanks Dr Daryl Lloyd at the Department for Transport for supplying the CO₂ estimate data and answering numerous queries about their calculation! The author also thanks the anonymous reviewers for their useful comments and suggestions. This work was funded by a 2012 Royal Geographical Society Small Grant and an Economic and Social Research Council research grant (ES/K007459/1).